

Homework #5 – Caches and Memory

Due Friday, April 12, at 5:00pm

You must do all work individually, and you must submit your work electronically via Sakai.

- For short-answer questions, submit your answers in PDF format as a file called <NetID>-hw5.pdf. **Word documents will not be accepted.**
- For the programming question, submit your source file using the filename specified in the question.

All submitted programs will be tested for suspicious similarities to other programs (from this semester and recent semesters), and the test will uncover cheating, even if it is “hidden.”

1) [5 points] Why are write-through caches usually also write-no-allocate?

2) [5] Your L1 data cache has an access latency of 3ns, and your L2 cache has an access latency of 20ns. Assume that 90% of your L1 accesses are hits, and assume that 100% of your L2 accesses are hits. What is the average memory latency as seen by the processor core? Show your work.

3) [25] You have a 64-bit machine and you bought 4GB of physical memory. Pages are 64KB. Remember that $1K=2^{10}$, $1M=2^{20}$, and $1G=2^{30}$.

[2] (a) How many virtual pages do you have per process?

[2] (b) How many physical pages do you have?

[2] (c) In the translation from a virtual address to a physical address, how many bits of VPN are you mapping to how many bits of PPN?

[2] (d) How big does a page table entry (PTE) need to be to hold just a single PPN?

[2] (e) How many PTEs fit on a page, assuming PTEs are the size computed in part (d)? (Real PTEs need to be somewhat bigger than this, but don't worry about that detail here.)

[2] (f) How many pointers fit on a page?

[2] (g) How big would a flat page table be for a single process, assuming PTEs are the size computed in

part (d)?

[3] (h) What are the virtual page offset bits for virtual address 25012? What are the physical page offset bits for virtual address 25012 after it has been translated?

[5] (i) Does a TLB miss always lead to a page fault? Why or why not?

4) [100 points] In Java *[or in C for up to 10 extra credit points]*, write a simulator of a cache.

The simulator, called `cachesim`, takes the following input parameters on the command line:

- name of the file holding the loads and stores
- cache configuration: size (not including tags or valid bits) in KB, associativity, and whether the cache is writeback and write-allocate (wb) or writethrough and write-no-allocate (wt)
- block size in bytes

The replacement policy is always LRU.

For example, `./cachesim tracefile 1024 4 wt 32` should simulate a cache that is 1024KB (=1MB), 4-way set-associative, write-through/write-no-allocate, with 32B blocks. Note, for java, to compile run `javac cachesim.java #to compile` and to simulate run `java cachesim traces/tracefile.txt 1024 4 wt 32`.

This cache hierarchy will be processing the loads and stores in the file called `tracefile`.

Important Assumptions: Addresses are 16-bits (2 bytes), and thus addresses range from 0 to $2^{16}-1$ (i.e., there is 64KB of address space). The machine is byte-addressed and big-endian. The cache size, associativity, block size, and access size will all be powers of 2. Cache size will be no larger than 16KB, block size will be no larger than 64B, and no access will be larger than the block size. No cache access will span multiple blocks (i.e., each cache access fits within a single block).

All cache blocks are initially invalid. All cache misses are satisfied by the main memory (and you must track the values written through to memory in case they are subsequently loaded). If a block has never been written before, then its value in main memory is zero.

The trace file will be in the following format. There will be some number of lines. Each line will specify a single load or store, the 16-bit address that is being accessed (in base-16), the size of the access in bytes, and the value to be written if the access is a store (in base-16). For example:

store 25bb 2 c77e

load d531 4

store 0000 1 01

Your simulator must produce the following output. For every access, it must print out what kind of access it is (load or store), what address it's accessing (in base-16, 4 digits), and what happens

(either “hit” or “miss”). For each load, it must also print out the value (in base-16, $2*N$ digits, N is the number of access size) that is loaded (possibly after satisfying the miss from memory).

The output format must be as follows (for the auto-grader to accept it):

store 25bb miss

load d5e1 hit 7d2f13ac

store 0000 miss

Running the tester: Use the following commands to run the provided tester:

- For a C program:
~/hw5_test_kit \$ make #compile
~/hw5_test_kit \$./hw5test_c.py ALL #run test
- For a Java Program:
~/hw5_test_kit \$ javac cachesim.java #compile
~/hw5_test_kit \$./hw5test_java.py ALL #run test

Note for implementing in C: You DO NOT need to worry about memory leaks for this assignment.

When you submit your program, if you have any known remaining bugs, please also include those in a README file to help the grader give you partial credit.